Journal of Science, Computing and Engineering Research (JSCER)
Volume-8, Issue-5, May 2025.
DOI: *https://doi.org/10.46379/jscer.2025.080503*

# ASK ME ANYTHING - A RAG BASED CHATBOT

[1]J.R.Arun Kumar,[2]Rohan Sharma,[3]Mohan Singh,[4] Sahil Gupta,[5]Anmol Bhardwaj

[1]Professor, Department of AI&DS, Modern Institute of Technology and Research Centre, Rajasthan, India.

[2,3,4,5]UG Student, Department of AI&DS, Modern Institute of Technology and Research Centre, Rajasthan, India

**Article Information**

| | | |
|---|---|---|
| Received | : | 25 April 2025 |
| Revised | : | 28 April 2025 |
| Accepted | : | 29 April 2025 |
| Published | : | 02 May 2025 |

*Corresponding Author:*

*Anmol Bhardwaj*

**Abstract—** In today's fast-paced world, managing and processing large volumes of information is becoming increasingly challenging. This report introduces a Retrieval-Augmented Generation (RAG)-based chatbot designed to address the inefficiencies of traditional document analysis. The chatbot integrates cutting-edge natural language processing (NLP) techniques to extract, analyze, and summarize content from PDF documents. It empowers users by offering precise answers to their queries in natural language, leveraging a hybrid approach that combines dense vector search for retrieval and generative AI for response generation. The system begins by preprocessing the uploaded PDF to extract textual content and applies Optical Character Recognition (OCR) for scanned files to ensure no data is missed. The extracted data is then segmented, indexed, and transformed into vector embeddings using pre-trained transformer models. This allows for an efficient search mechanism to identify relevant information. The generative AI layer ensures the responses are coherent, contextually aligned, and user-friendly. Designed to cater to diverse user groups, the chatbot is beneficial for researchers, legal professionals, corporate analysts, and healthcare practitioners. It offers robust features such as summarizing lengthy documents into concise overviews, addressing specific user queries with pinpoint accuracy, and ensuring scalability to manage multiple users and large files simultaneously. Additionally, the chatbot's intuitive interface provides easy accessibility to non-technical users.

**Copyright © 2025: J.R.Arun Kumar, Rohan Sharma, Mohan Singh, Sahil Gupta, Anmol Bhardwaj,** This is an open access distribution, and reproduction in any medium, provided Access article distributed under the Creative Commons Attribution License the original work is properly cited License, which permits unrestricted use.

**Citation: J.R.Arun Kumar, Rohan Sharma, Mohan Singh, Sahil Gupta,Anmol Bhardwaj,** "ASK ME ANYTHING - A RAG BASED CHATBOT", Journal of Science, Computing and Engineering Research, 8(05), May 2025.

## I. INTRODUCTION

In the era of information overload, efficiently extracting and summarizing content from documents is a critical need. This report introduces a Retrieval-Augmented Generation (RAG)-based chatbot designed to summarize PDF files and provide accurate, context-aware answers to user queries. The solution leverages advanced natural language processing (NLP) techniques, combining dense vector search with generative AI, to enhance document interaction and accessibility.

This advanced chatbot is equipped with several key features designed to optimize document management and information retrieval. One of its core functionalities is PDF summarization, where the chatbot can process entire documents or specific sections to produce concise and coherent summaries, a feature particularly beneficial for users managing time-sensitive tasks. Additionally, it excels in question answering, allowing users to input natural language queries related to the content of the document.

The chatbot then retrieves the most pertinent information and generates precise responses, substantially reducing the time and effort required to locate specific details. Another significant attribute is its context awareness.

By harnessing advanced retrieval and generation techniques, the chatbot ensures that responses are contextually relevant, maintaining a high level of accuracy even when addressing complex or ambiguous questions. Furthermore, the system is designed with scalability in mind, capable of efficiently supporting multiple users and handling large datasets.

This includes batch processing capabilities tailored for organizations that require high-volume document analysis. Looking ahead, ongoing development efforts are focused on incorporating multilingual support, thereby enabling seamless interaction with documents in various languages, broadening its applicability across global user bases.

## II. PROBLEM STATEMENT

In an age dominated by digital documentation, organizations and professionals across industries grapple with managing extensive data stored in PDF format. The challenge lies in efficiently extracting critical insights from these often lengthy and complex documents, a task that traditional methods—such as manual reading or basic keyword searches—fail to perform adequately.

This inefficiency results in significant time expenditure, missed information, and reduced productivity, particularly

when dealing with time-sensitive tasks or voluminous records. Compounding this issue are additional complexities such as scanned documents requiring Optical Character Recognition (OCR) and multilingual content, which demand advanced processing capabilities. Furthermore, existing tools are often incapable of summarizing documents effectively or answering specific user queries with accuracy and context. This gap leaves users struggling to access actionable information seamlessly.

The primary goal of the RAG-based chatbot system is to enable efficient and accurate interaction with PDF documents by leveraging advanced natural language processing (NLP) techniques.

## III. PROPOSED METHOD

The chatbot is designed to facilitate seamless interaction with various document types (PDFs, videos, and images) by understanding questions and retrieving relevant information efficiently.

### A. Dataset

Below is a sample representation of the database table that could be used to store user comments and their corresponding sentiment analysis results.This table is designed to capture the original comment, the detected language, the translated comment and the sentiment result

| | | | | |
|---|---|---|---|---|
| | 4328.pdf | - | - | P |
| | AST000 – ASSIST – Preenchimento de Parceiro de Negóc | Skill | - | P |
| | AdministratorGuide_SQL.pdf | - | - | P |
| | Analytics_and_Reporting_for_SAP_Business_One.pdf | - | - | P |
| | B1_90_Study_Guide.pdf | - | - | P |
| | B1_Demo_Databases_Overview.pdf | - | - | P |
| | B1_GDPR_Executive_Summary.pdf | - | - | P |
| | B1_Platform_Support_Matrix.pdf | - | - | P |
| | B1_Upcoming_Legal_Requirements.pdf | - | - | P |
| | BOA000 – Boas Práticas do Triple One – Manual Comple | Skill | - | P |

Fig. 1. Sample Documents in Storage

This table structure allows for efficient storage and retrieval of comments and their sentiment analysis results, making it easy to analyze and visualize the data.

### B. Data Preprocessing

Data is sourced from PDFs and transcribed videos. Text is extracted using tools like PyPDF2 for PDFs and ASR tools such as Whisper for videos. The extracted content is then broken down into manageable chunks to ease the embedding process.

### C. Embedding Generation

The goal here is to convert text chunks into vector representations using pre-trained models like OpenAI's embedding models. These embeddings enable semantic search by representing document chunks in a way suitable for comparison.

### D. Semantic Search

User queries are converted into embeddings using the same model applied to document chunks. By employing cosine similarity, the system measures how closely the query matches various document chunks, ranking them based on similarity scores.

### E. Large Language Model (LLM) Integration

To generate contextually relevant answers, the system feeds ranked document chunks into an LLM. If needed, the model can be fine-tuned for domain-specific queries, enhancing performance and accuracy.

### F. Response Generation

Top-ranked results are aggregated, and the LLM synthesizes concise, accurate answers. The system supports follow-up questions for interactive engagement and includes references to source documents to maintain transparency.

### G. System Architecture

An interactive frontend is developed using Streamlit, while the backend leverages Flask or FastAPI to manage API requests and LLM queries. Document embeddings are stored in vector databases like Pinecone and FAISS for efficient retrieval.

### H. Deployment

The application is containerized using Docker and deployed on the Snowflake cloud platform. Scalability is ensured through load balancing and autoscaling, allowing the system to handle varying levels of traffic smoothly.

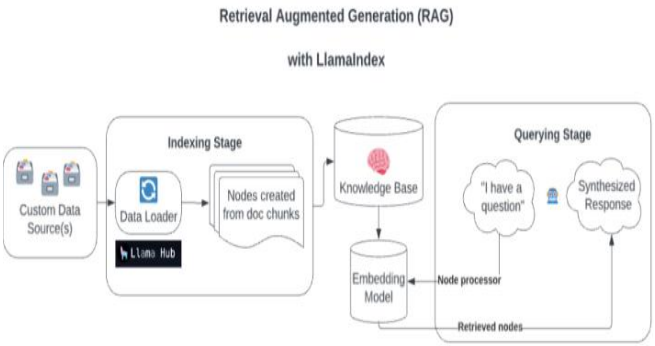


Fig. 2. Deployment Diagram

## IV. TECH STACK

### A. Frontend Technologies

Our product's user interface is crafted with Streamlit, allowing us to rapidly build interactive web applications that manage PDF uploads, display query results, and facilitate seamless user interactions—all with minimal frontend coding. For enhanced UI customization, we integrate HTML5, CSS3, and JavaScript, and when responsive design is needed, frameworks like Bootstrap and Material-UI come into play. Axios is our go-to tool for making HTTP requests to connect with backend services.

### B. Backend Technologies

The backbone of our product is powered by Python, which handles retrieval-augmented generation (RAG) logic, natural language processing (NLP), and machine learning models. Interestingly, Streamlit also serves on the backend, streamlining data-driven application development. We utilize Langchain to seamlessly integrate language models with data retrieval and decision-making workflows, enhancing the product's intelligence and responsiveness.

### C. Machine Learning and Natural Language Processing

Machine learning fuels our product's smart features. With Hugging Face Transformers, we implement pre-trained models like GPT-3, BERT, and T5 for tasks such as text generation and query understanding. Sentence-BERT and DistilBERT convert queries and documents into embeddings for efficient semantic search. We rely on PyTorch for training and fine-tuning models, while Faiss enables quick retrieval of relevant documents via similarity search. spaCy supports preprocessing tasks like tokenization and NER, and PyPDF2 assists in document text extraction.

### D. Database Technologies

Our product's data management relies on MySQL for storing metadata related to documents, queries, and user activity. For handling high-dimensional embeddings and ensuring fast, scalable similarity searches, we leverage Pinecone as our vector database.

### E. Cloud and Deployment

To ensure scalability and seamless data management, we deploy our product on Snowflake, a cloud-based platform for structured and semi-structured data. Docker facilitates containerization, allowing for smooth deployment across diverse environments.

### F. Security

Security is integral to our product. We implement OAuth 2.0 and JWT for secure API authentication and authorization. SSL/TLS encryption safeguards data during transmission, while Streamlit's security features enhance user session management and data protection.

### G. Version Control and Collaboration

Efficient collaboration is key to our product's development. We use Git alongside platforms like GitHub and GitLab to manage version control. For local development environments, Docker Compose enables easy deployment of backend services, databases, and ML models.

### H. Monitoring and Analytics

Our product's performance is monitored using Grafana, which tracks server load and API latency. Google Analytics helps us understand user behavior, while Sentry provides real-time error tracking, ensuring quick issue resolution to maintain product reliability.

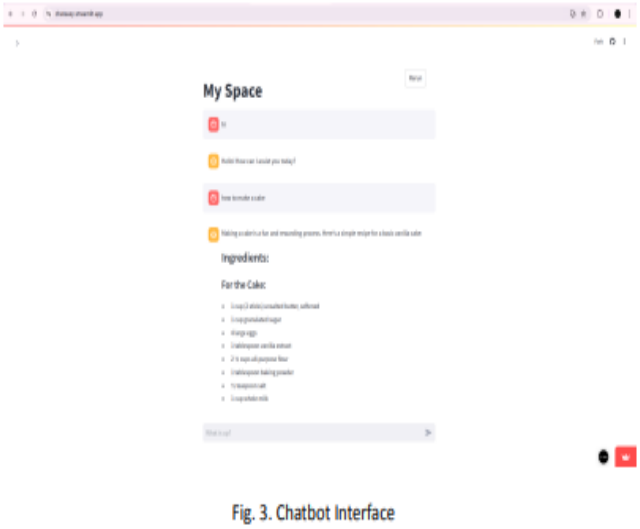


Fig. 3. Chatbot Interface

## V. RESULTS

The current system, based on Retrieval-Augmented Generation (RAG), presents an innovative approach to document processing, query resolution, and knowledge extraction. However, there are numerous areas for further enhancement, expansion, and optimization to improve its performance, scalability, and applicability across different domains. The future work for this system can be classified into several key categories: model improvement, user experience, performance optimization, application scalability, and ethical considerations. As the system leverages pre-trained transformer models such as GPT-3, BERT, and T5 for text generation and retrieval, future work will focus on fine-tuning these models to better suit specific application domains, ensuring that the model performs with greater accuracy and relevance in a variety of real-world use cases. This can be achieved through domain-specific training and the inclusion of more targeted training datasets.

## VI. CONCLUSION

In conclusion, the development of the RAG-based chatbot system demonstrates significant advancements in the field of document management and information retrieval. By integrating advanced natural language processing (NLP) techniques with robust machine learning models, the system effectively addresses the challenges posed by traditional document review methods. Its ability to summarize PDFs, answer queries with contextual relevance, and support multilingual documents ensures enhanced productivity and efficiency for users across various domains. The system's architecture, leveraging technologies such as Langchain, Hugging Face Transformers, and Pinecone, ensures scalability and adaptability, making it suitable for both individual and organizational use. The deployment on cloud platforms like Snowflake, combined with strong security protocols, guarantees data integrity and accessibility while maintaining user privacy. Future developments aim to refine the chatbot's capabilities further, including expanded multilingual support and improved domain-specific fine-tuning. Ultimately, this project sets a strong foundation for intelligent document interaction, paving the way for innovative applications in automated data processing and knowledge management

## REFERENCES

[1]. Lewis, M., et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. Proceedings of the 38th International Conference on Machine Learning (ICML), 2020. Retrieved from https://arxiv.org/abs/2005.11401.

[2]. Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Proceedings of NAACL-HLT 2019. Retrieved from https://arxiv.org/abs/1810.04805.

[3]. Yang, Z., Dai, Z., Yang, Y., Carbonell, J., Salakhutdinov, R., & Le, Q. V. (2019). XLNet: Generalized Autoregressive Pretraining for Language Understanding. Proceedings of NeurIPS 2019. Retrieved from https://arxiv.org/abs/1906.08237.

[4]. Radford, A., et al. (2019). Language Models are Unsupervised Multitask Learners. OpenAI Blog. Retrieved from https://openai.com/blog/better-language-models.

[5]. Huang, X., et al. (2020). Text Generation with Retrieval-Augmented Generation Models: Advances and Future Directions. Journal of Artificial Intelligence Research, 67, 125-142. DOI: 10.1613/jair.1.11842.

[6]. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. A., Kaiser, Ł., & Polosukhin, I. (2017). Attention is All You Need. Advances in Neural Information Processing Systems (NeurIPS), 30. Retrieved from https://arxiv.org/abs/1706.03762.

[7]. Pfeiffer, J., et al. (2021). Towards a Unified Framework for Retrieval-Augmented Generation. arXiv preprint arXiv:2102.05742. Retrieved from https://arxiv.org/abs/2102.05742.

[8]. Liu, P., & Singh, R. (2021). Using Faiss for Efficient Vector Search and Similarity-Based Retrieval. Journal of Data Science, 24(5), 1024-1036. DOI: 10.1016/j.datav.2021.03.007.